

The SnuPL/1 Intel IA-32 Backend

This document gives a minimal introduction to the Intel IA-32 instruction set and calling convention. Intel provides detailed manuals at the following address:

<http://www.intel.com/content/www/us/en/processors/architectures-software-developer-manuals.html>

[Volume 2](#) contains the complete instruction set reference of Intel IA-32 processors.

The Application Binary Interface that defines the calling convention and interaction with the operating system/runtime environment is described in [System V ABI for Intel386 Architectures](#).

Refer to the The SnuPL/1 Language for links to x86_64 (AMD64) processors.

1. Registers

IA-32 has 8 general-purpose registers, six of which can be used more or less arbitrarily.

`eax, ebx, ecx, edx, esi, edi, ebp, esp`

esp/ebp represent the stack pointer and the base pointer; both registers are used to implement the calling convention. The other registers can be used freely with few exceptions (for example, multiplication and division use predefined registers). IA-32 is fully backwards compatible; this is why the registers can be accessed in their old 16- or 8-bit form. This is no concern for us, we will only use the full 32-bit registers.

Not visible here are two important registers: the program counter and the condition codes. The program counter is manipulated indirectly through control-flow instructions, and the condition codes are set/read implicitly by ALU operations and conditional branches.

2. Instructions

The instructions listed in the table on the following page suffice to implement a simple code generator for SnuPL/1. We use GCC to assemble our programs, hence the assembler syntax below uses the AT&T syntax. In AT&T syntax, the source is listed *before* the destination, immediate values are prefixed with a "\$", and registers are prefixed with a "%" character.

Memory addresses have the form

`displacement(%base, %index, scaling factor)`

and the accessed location is

`mem[%base+ %index * scaling factor + displacement]`

We only require two sub-forms: `disp` and `disp(%base)` to access globals and locals/parameters/temps, respectively.


Instruction	Effect	Description
addl S, D	$D \leftarrow D + S$	addition
subl S, D	$D \leftarrow D - S$	subtraction
andl S, D	$D \leftarrow D \&\& S$	logical and
orl S, D	$D \leftarrow D \parallel S$	logical or
negl D	$D \leftarrow -D$	negate
notl D	$D \leftarrow \sim D$	logical not
imull S	$[EDX:EAX] \leftarrow [EAX] * S$	32-bit signed multiply
idivl S	$[EAX] \leftarrow [EDX:EAX] / S$	32-bit signed division
cmpl S2, S1	$[\text{condition codes}] \leftarrow S1 - S2$	set condition codes based on the comparison S1, S2
movl S, D	$D \leftarrow S$	move
cdq	$[EDX:EAX] \leftarrow \text{sign_extend}([EAX])$	sign-extend to 64-bit
pushl S	$[ESP] \leftarrow [ESP] - 4$ $\text{mem}[ESP] \leftarrow S$	push S onto stack
popl D	$D \leftarrow \text{mem}[ESP]$ $[ESP] \leftarrow [ESP] + 4$	pop top of stack into D
call T	push return address continue execution at T	subroutine call
ret	pop return address from stack continue execution at return address	subroutine return
jmp T	continue execution at T	unconditional branch
je T	goto T if condition codes signal "equal"	conditional branch
jne T	goto T if condition codes signal "not equal"	
jl T	goto T if condition codes signal "less than"	
jle T	goto T if condition codes signal "less equal"	
jg T	goto T if condition codes signal "bigger than"	
jge T	goto T if condition codes signal "bigger equal"	
nop		no operation

For almost all arithmetic instructions, one of the operands (but not both) can be a memory address (a notable exception is the `idivl` instruction), a register, or an immediate. The other operand is an immediate or a register.

3. Data: parameters, local variables and temporaries are stored on the stack and addressed relative to the stack and/or base pointer. Global data, however, must be allocated statically. The assembler allows to give names (labels) to junks of data, and you can then use those names directly as operands of instructions. Use the `.long <val>` and `.skip n` assembly directives to allocate an initialized long value or *n* uninitialized bytes of memory, respectively. Do not forget to initialize the meta-data of arrays (dimensions) and the content of string constants.

4. Calling convention: in the Intel IA-32 calling convention, procedure activation frames are built by the stack and base pointer. The stack grows towards smaller addresses. The stack pointer points to the top of the stack. Storing data below the stack pointer is not allowed. Function arguments to subroutine calls are passed on the stack in reverse order, i.e., the first argument is on top of the stack, the second one immediately below the first one, and so on. Function return values, if present, are returned in register `eax`. The registers `ebx`, `esi`, and `edi` are callee-saved and thus must be preserved across function calls. Implicitly, `esp` and `ebp` are also callee-saved (both `esp` and `ebp` must point to the caller's activation record after returning from a subroutine call).

5. Procedure/function activation frame: below we give one possibility of a procedure activation frame. You are free to choose your own layout.



In the design above, procedure/function parameters are pushed onto the stack immediately before the call (and must be removed after returning). This is more convenient than a pre-computed fixed parameter area when supporting nested function calls.

The parameters and the return address are generated by the caller. The parameters by a series of `push` instructions, the return address implicitly by the `call` instruction. Upon entering a function, the callee has to create the remaining parts of the activation frame as follows:

1. save `ebp` by pushing it onto the stack
2. set `ebp` to `esp`
3. save callee-saved registers
4. generated space on the stack for locals and spilled variables by adjusting the stack pointer

Immediately before returning to the caller, the callee needs to restore the callee-saved registers and dismantle the activation frame. This can be achieved by the following steps

1. remove space on stack for locals and spilled variables by setting the stack pointer immediately below the callee-saved registers.
2. restore callee-saved registers
3. restore `ebp`
4. issue the `ret` instruction

After returning from the callee, the caller has to remove the parameters from the stack.

Appendix 1: AT&T Assembly, Assembling and Linking with the I/O Routines

Assembly programs are structured into sections. For our purposes, we require two sections: the `.text` section contains assembled machine code, while the `.data` section holds global variables.

Here is a skeleton file for programs in AT&T syntax:

```
# template

.text                # beginning of the text section
.align 4             # align text section at a 4-byte boundary

.global main         # to let the assembler know that we implement
                    # the function "main" (= module body)

.extern ReadInt       # externally defined functions (I/O, array handling)
...
main:                # module body, followed by functions/procedures
...

.data                # beginning of the data section
.align 4             # align at a 4-byte boundary

p: .long 1           # global array 'p': integer[10]
   .long 4
   .skip 40
x: .skip 1           # global variable 'x' (1 byte)

.end                # end of program
```

Be aware that labels must be local or unique. An easy way of generating unique labels is to prefix them with the name of the scope (i.e., the procedure name) they are defined in.

The assembly file generated by `snuplc` can be compiled using `gcc` as follows

```
$ gcc -m32 -o primes.o -c primes.mod.s
```

The `-m32` options tells GCC that we want to build a 32-bit binary, and `-c` instructs the assembler just to assemble the input file into an object file.

To generate an executable file, the object file is linked together with the provided array and I/O routines as follows (here we assume that the libraries are located in subdirectory `rte/IA32/`)

```
$ gcc -m32 -o primes primes.o rte/IA32/ARRAY.s rte/IA32/IO.s
```

Of course, this can also be done in one step:

```
$ gcc -m32 -o primes primes.mod.s rte/IA32/ARRAY.s rte/IA32/IO.s
```

The SnuPL/1 compiler can execute these commands for you if provided with the `--exe` option

```
$ snuplc --exe primes.mod
```

Now you can run your program with

```
$ ./primes
```

Appendix 2: GNU Debugger (GDB)

Debugging a generated assembly file can be challenging without a nice IDE. You may need to debug your x86 program using a command line debugger and follow the execution instruction-by-instruction to see what's going on.

The GNU debugger, `gdb`, is an excellent tool to debug programs. While it seems to be rather crude, it offers lots and lots of functions that many people are not aware of.

To start debugging your program using `gdb`, type

\$ `gdb ./primes`

at the prompt. GDB greets you with

```
GNU gdb (Gentoo 7.12.1 vanilla) 7.12.1
Copyright (C) 2017 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.  Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-pc-linux-gnu".
Type "show configuration" for configuration details.
...
For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from ./primes...(no debugging symbols found)...done.
(gdb)
```

From there, commands will help you run/stop/break your program and inspect/modify values. The following table contains a list of handy commands that you might use when debugging your program. For a complete list, type `help` and follow the instructions on the screen.

Command	Description
<code>r (un)</code>	run the program until - it ends - it crashes - it hits a breakpoint
<code>c (ontinue)</code>	continue a stopped program
<code>quit</code>	exit GDB
<code>break *address</code>	set a breakpoint at address
<code>break name</code>	set a breakpoint at name
<code>stepi</code>	execute one assembly instruction
<code>stepi n</code>	execute n assembly instructions
<code>disas</code>	disassemble around current program counter
<code>disas name</code>	disassemble at name
<code>display /5i \$pc</code>	disassemble the next 5 instructions at <code>pc</code> at every stop
<code>display \$<reg></code>	display the value of <code>reg</code> at every stop
<code><Enter></code>	re-run the last command

Especially the `display` command together with `stepi` will be very helpful when stepping through your program.